

# Corbis for Academic Research


Getting started with Corbis and the Literature Starter Kit

## Who this is for

Researchers who want to use AI assistants for literature reviews, idea screening, citation management, and a disciplined search trail.

**Setup:** Claude Code, Codex, or Cursor with Corbis MCP.

**Goal:** Leave with a first workflow and first prompts to run.



# What Corbis Is

## Research-grade AI tools

Corbis gives your AI assistant access to a curated academic corpus and returns results with citations you can inspect and verify.

Start with `search_papers`, `get_paper_details`, `top_cited_articles`, and `export_citations`.

## What it provides via MCP

- Paper search, metadata, and citation export
- FRED macro data
- CRE market intelligence and dataset discovery
- Web search and deep research (select plans)

## Why researchers should care

- Search the literature before making claims
- Trace where every piece of evidence came from
- Move from search to review to figures to writing in one repo

# Why This Repository Exists

## The problem

AI assistants are fluent but undisciplined: they skip searches, hide provenance, and leave no research trail.

## What this kit adds

- 6 reusable research workflows
- Structured outputs in notes/ and output/
- Automatic BibTeX export and LaTeX handoff
- Transparent search logging

## What you gain

- Map a field in minutes, not days
- Screen ideas before committing weeks of work
- Keep a clean citation and writing pipeline
- Maintain a reproducible record of every search

# What a New User Needs

## Required

- A Corbis account
- OAuth sign-in to Corbis
- An MCP-capable AI client

## Recommended

- Claude Code for the smoothest slash-command flow
- Basic comfort with repo files
- Check citations before using them

## Optional

- Python 3.10+ for figure generation
- LaTeX for manuscript drafting
- An MCP API key for key-based authentication
- Codex or Cursor as alternative clients

## Beginner rule

Always search before asserting novelty, importance, or data availability.

# First-Time Setup in Claude Code

## Step 1: Use your Corbis account

OAuth sign-in works without an MCP API key. The repository's MCP configuration contains only the Corbis server URL.

## Step 2: Connect Claude Code

```
cd corbis-literature-starter-kit
claude
```

Complete the OAuth sign-in when prompted.

## Step 3: Work inside this repo

Clone the starter kit, open it in Claude Code, and run your first workflow: `/lit-review [your topic]`.

## Optional key-based setup

If your client cannot use OAuth, create an MCP API key in Corbis Settings and configure a private Authorization header. Never place a key in a URL or a committed file.

# Other Client Paths

## Codex

- Add Corbis to your `config.toml` (global or project-level)
- Add the MCP server URL, then run `codex mcp login corbis`
- Verify the connection with `codex mcp list`
- Full guide: `CORBIS_MCP_CODEX_GUIDE.md`

## Cursor

- Open **Settings > MCP Servers > Add**
- Name: `corbis`
- Paste the MCP endpoint URL and complete OAuth sign-in
- Reload, then use workflow names as prompts
- Full guide: `CORBIS_CURSOR_PLUGIN.md`

## All clients, same tools

Claude Code has the smoothest slash-command experience. Codex and Cursor use the same Corbis tools and the same workflow prompts from this repo.

# Repository Map

| Path                           | Why a beginner should care                                                     |
|--------------------------------|--------------------------------------------------------------------------------|
| <code>.claude/commands/</code> | Slash-command wrappers ( <code>/lit-review</code> , <code>/idea</code> , etc.) |
| <code>.claude/skills/</code>   | Workflow definitions behind each command                                       |
| <code>.claude/agents/</code>   | Paper-reader prompt for deep reading of key PDFs                               |
| <code>notes/</code>            | Lab notebook and dated research decisions                                      |
| <code>output/</code>           | Reviews, idea cards, figures, <code>paper_set.json</code> , search log         |
| <code>latex_template/</code>   | Clean article template ( <code>natbib</code> + <code>plainnat</code> )         |
| <code>utils/</code>            | Python figure generator for <code>/lit-landscape</code>                        |
| <code>references/</code>       | Writing norms and citation formatting                                          |

# How the Workflows Connect

## Shared files

- `output/paper_set.json`: the canonical paper dataset reused across workflows
- `output/search_log.md`: a transparent record of what the assistant actually searched

## The key idea

Run one workflow to build a literature base. Later workflows reuse it instead of searching from scratch.

## Cross-client use

In Claude Code these are slash commands. In Codex or Cursor, use the same names and examples as prompt templates.

## Typical chain

- 1 `/lit-review [topic]` builds the paper set.
- 2 `/lit-landscape [topic]` turns that set into figures.
- 3 `/brainstorm [topic]` proposes ideas against the same literature.
- 4 `/idea [specific idea]` stress-tests one candidate.

# Skill 1: /lit-review

## Use it when

You need to understand a field before pitching, writing, or designing a project.

## Produces

- A synthesized review organized by theme
- A reading list of central papers
- BibTeX citations ready for your manuscript
- A reusable paper set for later workflows

## Try it

`/lit-review climate risk and commercial real estate`

- Good first move for a seminar paper or dissertation scan
- Clusters papers by theme, not just by keyword
- Leaves you with a reusable paper set for the next workflow

## Skill 2: `/lit-search`

### Use it when

You already have a question or paper idea and need to locate the closest literature before claiming a contribution.

### Produces

- Closest 3–5 papers
- A draft contribution paragraph
- A related-literature outline
- BibTeX export for cited papers

### Try it

`/lit-search` Do bank branch closures reduce small business lending?

- Useful before a proposal meeting or referee response
- Searches by concentric rings: direct competitors, nearby methods, seminal work, recent papers
- Forces you to articulate how your paper differs

## Skill 3: /brainstorm

### Use it when

You know the broad area but not which question is worth pursuing.

### Produces

- 10 ranked ideas from a wider candidate pool
- Expanded sketches for the top 3
- A recommendation for which idea to screen next

### Try it

/brainstorm behavioral biases in household mortgage decisions

- Good for early-stage dissertation exploration
- Applies multiple idea-generation lenses, then checks novelty against Corbis
- Kills weak ideas before you become attached to them

## Skill 4: /idea

### Use it when

You have one concrete idea and want a hard-nosed screen, not encouragement.

### Produces

- An Idea Card with scoring and closest papers
- Novelty verification against Corbis
- Stress tests and essential hurdles
- A Go, Revise, or Kill verdict

### Try it

/idea Do zoning reforms lower rent growth by changing landlord expectations?

- Useful before spending weeks on data work
- Runs a desk-editor screen and three stress tests
- Identifies the hurdles you must clear and the exhibit you need

## Skill 5: `/verify-citations`

### Use it when

You have a LaTeX manuscript and want to audit whether every citation is real, complete, and correctly matched.

### Produces

- A verification report
- Per-citation labels: OK, PARTIAL, MISMATCH, UNVERIFIED, or MISSING
- Checks against both the `.tex` and `.bib` files

### Try it

`/verify-citations`

- Final pass before circulating or submitting
- Catches stale BibTeX entries copied across projects
- Reduces embarrassing citation errors

## Skill 6: /lit-landscape

### Use it when

You want figures that show trends, methods, and gaps in a literature.

### Produces

- Publication timelines and citation landmark charts
- Journal distributions and method evolution
- Thematic heatmaps and coverage gap maps

### Try it

`/lit-landscape corporate governance and firm performance`

- Useful for dissertations, review talks, or appendix figures
- Reuses `paper_set.json` when available
- Requires Python: `pip install -r requirements.txt`

# Recommended Workflow

- 1 `/lit-review [topic]` to learn the field and build the paper set
- 2 `/brainstorm [topic]` to generate ideas, or `/lit-search [question]` to position an existing one
- 3 `/idea [specific idea]` to screen before committing to heavy work
- 4 `/lit-landscape [topic]` for visual evidence (talks, prospectus, appendix)
- 5 Draft with `output/` files and the `latex_template/`
- 6 `/verify-citations` before circulating

## The discipline

Use Corbis to narrow the field. Use this repo to keep the search trail, outputs, and citations organized.

# Outputs and Research Hygiene

## What should stay in your repo

- `notes/lab_notebook.md` for dated research decisions
- `output/*.md` for reviews, memos, and idea cards
- `output/paper_set.json` and `output/search_log.md`
- `.bib` exports plus the LaTeX manuscript files

## Corbis extras outside the repo



Research Insights gives journal-level corpus context. Open Datasets helps pair literature ideas with empirical data sources.

## Why this matters

The repo gives you a search trail you can revisit, share, and reuse when the project evolves.

# First Prompts to Try

## Start with these three

- 1 `/lit-review climate risk and commercial real estate pricing`
- 2 `/brainstorm behavioral biases in household mortgage decisions`
- 3 `/idea Do bank branch closures reduce small business lending through relationship destruction?`

## Docs to keep nearby

- `README.md`
- `SKILLS_USE_GUIDE.md`
- `CORBIS_MCP_CODEX_GUIDE.md`
- `CORBIS_MCP_TOOL_REFERENCE.md`

## Remember

These workflows are not magic. They are a disciplined loop: search, synthesize, screen, document, verify.

`corbis.ai`    `www.corbis.ai/docs`